

Neural Network Regularization Data Augmentation and Dropout

CS109B Data Science 2

Pavlos Protopapas, Mark Glickman



Outline

Regularization of NN

- Norm Penalties
- Early Stopping
- **Data Augmentation**
- Dropout

Keys

Using the functional API

DataGenerator

Custom Loss

Custom Layer

Gradient Tape

TF Data

TF Records

When you move on to Deep Learning



Data Augmentation



hue



crop-and-pan



elastic



flip-lr




flip-ud



rotate




Data Augmentation: dos and don'ts




We use **ImageDataGenerator** to augment the dataset

```
def get_generator():  
    # create duplicate images  
    BATCHES_PER_EPOCH = 300//BATCH_SIZE  
    classes = ['pavlos', 'not-pavlos']  
    for img_class in classes:  
        img = Image.open(f'{DATA_DIR}/{img_class}.jpeg')  
        for i in range(1, BATCH_SIZE*BATCHES_PER_EPOCH//2+1):  
            img.thumbnail(TARGET_SIZE, Image.ANTIALIAS)  
            img.save(f'{DATA_DIR}/{img_class}/{img_class}{i:0>3}.jpeg', "JPEG")  
  
    data_gen = ImageDataGenerator(  
        rescale=1./255,  
        height_shift_range=0.5,  
        width_shift_range=0.5)  
  
    img_generator = data_gen.flow_from_directory(  
        DATA_DIR,  
        target_size=(TARGET_SIZE),  
        batch_size=BATCH_SIZE,  
        classes=classes,  
        class_mode='binary')  
    return img_generator
```



Data Augmentation: dos and don'ts

Carefully choose your transformations. Not all transformations are valid.



Data Augmentation does not work for tabular data and not as nicely for time series.

Outline

Regularization of NN



- Norm Penalties
- Early Stopping
- Data Augmentation
- **Dropout**

Co-adaptation

Overfitting occurs when the model is **sensitive** to slight variations on the input and therefore it fits the noise.

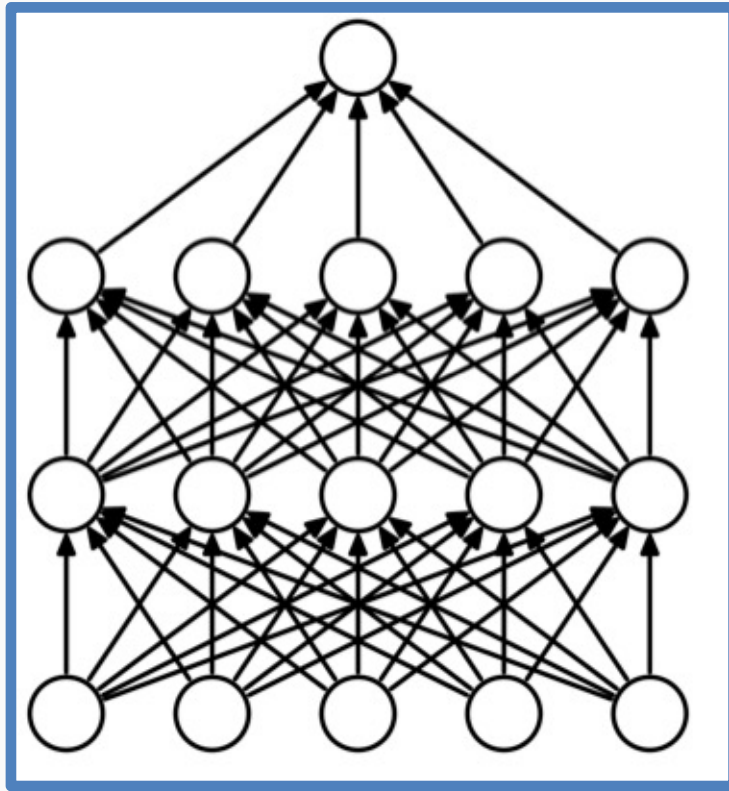
L1 and L2 regularizations ‘shrink’ the weights to avoid this problem.

However in a large network many units can collaborate to respond to the input while the weights can remain relatively small. This is called co-adaptation.




Dropout

- Randomly set some neurons and their connections to zero (i.e. “dropped”)
- Prevent overfitting by reducing **co-adaptation** of neurons
- Like training many random sub-networks



Standard Neural Network



After applying dropout

Dropout: Training

For each new example in a mini-batch (could be for one mini-batch depending on the implementation):

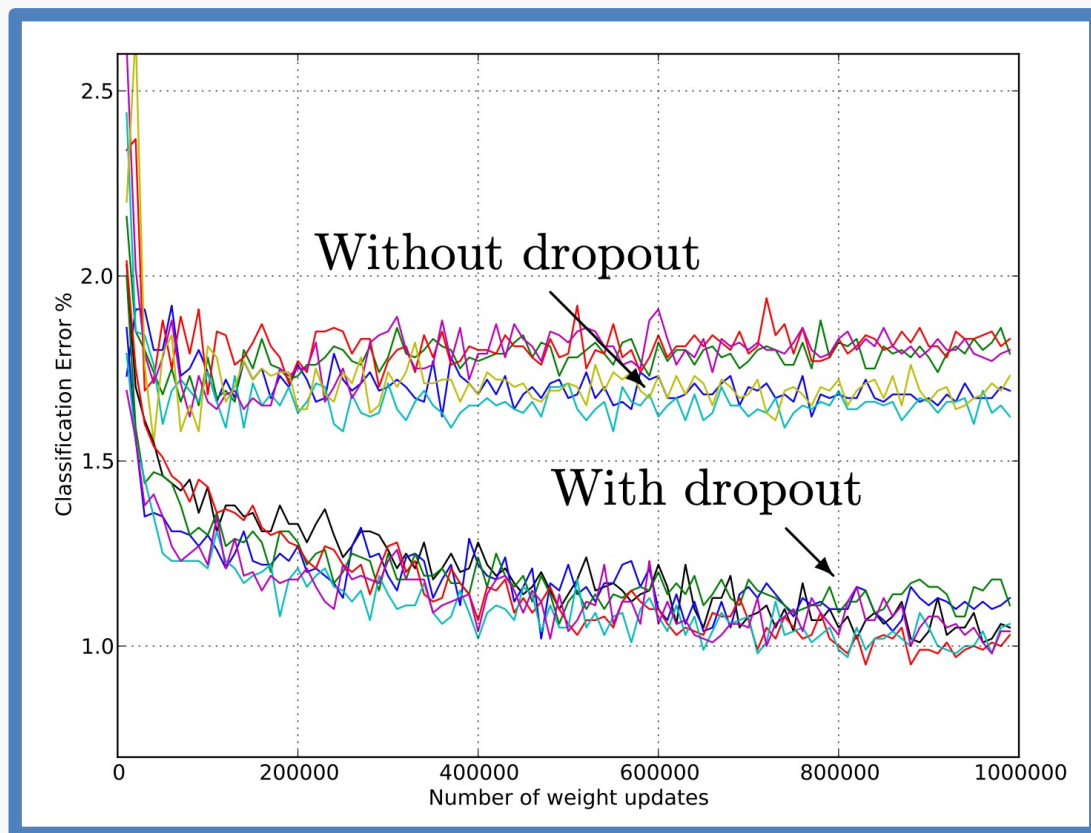
- Randomly **sample a binary mask** μ independently, where μ_i indicates if input/hidden node i is included
- **Multiply output of node i with μ_i** , and perform gradient update

Typically:

- Input nodes are included with prob=0.8 (as per original paper, but rarely used)
- Hidden nodes are included with prob=0.5

Dropout

- Widely used and highly effective





Test error for different architectures with and without dropout.

The networks have 2 to 4 hidden layers each with 1024 to 2048 units.

- Proposed as an alternative to ensemble methods, which is too expensive for neural nets

Dropout: Prediction

- We can think of dropout as training many of sub-networks
- At **test time**, we can “aggregate” over these sub-networks by **reducing connection weights in proportion to dropout probability, p**



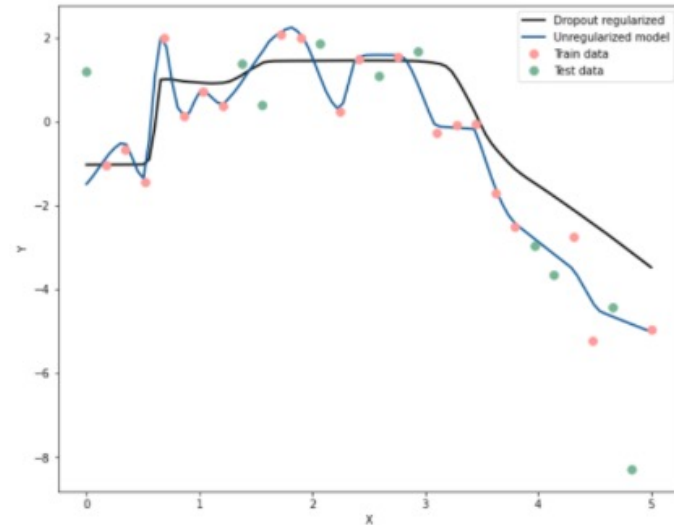
NOTE: Dropouts can be used for **neural network inference** by dropping during predictions and predicting multiple times to get a distribution

Exercise: Dropout

The goal of this exercise is to understand and use **dropouts** for neural network regularization.

This method avoids overfitting by briefly *switching off* certain weights during training.

NOTE: This graph is only a sample.



Instructions:

- Use the helper function `unregularized_model` to:
 - Generate the predictor and response data using the helper code given.
 - Build a simple neural network with 5 hidden layers with 100 neurons each and display the trace plot. This network has no regularization.
- For the same model architecture implement dropout by adding appropriate dropout layers.
- Compile the model with MSE as the loss. Fit the model on the training data.
- Use the helper code to visualise the MSE of the train and test data with respect to the epochs.
- Predict on the entire data.
- Use the helper code to plot the predictions along with the generated data.
- This plot will consist of the predictions of both the neural networks. The graph will look similar to the one given above.

